

Lesson 1 - Frequency and Response Latencies

Erin M. Buchanan

01/16/2019

Language Topics Discussed

- ▶ English Lexicon Project
- ▶ The Subtitle Projects
- ▶ Extension to Priming Projects

English Lexicon Project

- ▶ <http://ellexicon.wustl.edu/>
- ▶ ELP was a large undertaking that helped kick start the recent uptick in standardized behavioral data for language researchers
 - ▶ Corpora have been around, so have databases, but in the last 10 years, this field has grown exponentially
- ▶ The data contains 40K + words and 40K + nonwords with many characteristics
- ▶ Lexical Decision and Naming Tasks included

English Lexicon Project

- ▶ <https://www.psychtoolkit.org/experiment-library/ldt.html>
 - ▶ However, you would only see one word at a time (in this example you see two)
- ▶ In the naming task, you would be asked to read those words aloud, one at a time
- ▶ Output we are interested in is how the lexical variables might predict the response latencies

Subtitle Projects

- ▶ Starting with Brysbaert and New (several papers), there was a movement to rethink frequency and its relation to predicting language results
- ▶ Traditionally, two sources of frequency were used:
 - ▶ The Brown Corpus: Kucera and Francis (1967)
 - ▶ HAL Corpus: Burgess and Livesay (1998)
- ▶ However, we weren't sure these were the best estimators of frequency

Subtitle Projects

- ▶ <http://subtlexus.lexique.org/>
- ▶ Downloaded subtitles from www.opensubtitles.org with 50 million words+
- ▶ Provided both estimates for subtitles and music lyrics
- ▶ Models estimating lexical decision and naming times indicate these estimations of frequency are better predictors
- ▶ Expanded to 15-20 different languages

The Semantic Priming Project

- ▶ Priming occurs when cognitive processing is speeded because of a previous event
- ▶ Generally, we measure priming using lexical decision and naming tasks
- ▶ Let's say you have two trials:
 - ▶ DOCTOR → TREE (unrelated)
 - ▶ DOCTOR → NURSE (related)

The Semantic Priming Project

- Priming is thought to occur by several different mechanisms: spreading activation, deliberate cognitive processes such as expectancy generation, etc.



The Semantic Priming Project

- ▶ Contains lexical decision and naming response latencies for related, unrelated, and nonword trials
- ▶ Is paired with the ELP and SUBTLEX projects
- ▶ Gives us more data to predict either response latencies or priming latencies

Regression

- ▶ Simple regression is the relationship between one independent and one dependent variable (also correlation)
- ▶ Multiple regression is the relationship between two or more independent variables and one dependent variable
 - ▶ Useful because it allows you to examine the predictive ability of each variable adjusting for the other variables
- ▶ We can fit parametric (linear) models or nonparametric models, depending on the expectation of linearity, as well as the type of dependent variable


Understand Regression Models

$$\hat{y}_i = b_0 + b_1x_{1i} + b_2x_{2i}\dots + \epsilon_i$$

- ▶ \hat{y} is the predicted score for each person (i) on the dependent variable
- ▶ B_0 is the y-intercept
- ▶ B_1+ are the slope values for each predictor
 - ▶ Slopes are interpreted as *for every one unit increase in X , we see B unit increases in Y*
- ▶ X is each individual predictor
- ▶ Error for each individual person, as we never get their predicted score exactly right

Understand Regression Models

- ▶ Least Squares estimation
 - ▶ Creates the line of best fit by minimizing the residual error ϵ



Understand Regression Models

- ▶ For the overall model including all variables:
 - ▶ Determine statistical significance by using p values from an F -test for linear models
 - ▶ Determine practical significance by using R^2
- ▶ For the individual predictors:
 - ▶ Determine statistical significance by using p values from a t -test
 - ▶ Determine practical significance by using partial correlation pr^2

Examples Using ELP

- ▶ Word is the word presented to the participant
- ▶ Length is the number of characters in each word
- ▶ SUBTLWF is the subtitle word frequency estimate
- ▶ POS is part of speech
- ▶ Mean_RT is the mean response latency in milliseconds

```
library(Rling)
data(ELP)
head(ELP)
```

##	Word	Length	SUBTLWF	POS	Mean_RT
## 1	rackets	7	0.96	NN	790.87
## 2	stepmother	10	4.24	NN	692.55
## 3	delineated	10	0.04	VB	960.45
## 4	swimmers	8	1.49	NN	771.13
## 5	umpire	6	1.06	NN	882.50
## 6	cobra	5	3.33	NN	645.85

Dealing with Categorical Predictors

- ▶ How can we interpret and use categorical predictors?
- ▶ When X is continuous, the interpretation is that *for every one unit increase in X , we see B unit increases in Y*
- ▶ That doesn't work as well for categorical predictors. . .
- ▶ Instead, we have to use Dummy Coding (well, R does it for us)

Dummy Coding

- ▶ A way to represent categorical data for regression/least squares analyses
- ▶ You will get (categories - 1) predictors
- ▶ How to interpret these predictors?
 - ▶ Each predictor represents the difference in means between the coded group (the one with the 1) and the group coded as all zeroes (the “control” group)

Dummy V

No Affiliation

Indie Kid

Dummy Coding

- ▶ POS is a categorical predictor we want to use
- ▶ Three categories:
 - ▶ JJ: Adjective
 - ▶ NN: Noun
 - ▶ VB: Verb
- ▶ Default is to make the first category the comparison category

```
table(ELP$POS)
```

```
##
```

```
##   JJ   NN   VB
```

```
## 159 532 189
```

Dummy Coding

- Generally, nouns would be considered the comparison group, so let's rearrange them so they are "first".

```
ELP$POS = factor(ELP$POS, #the column you want to update  
                  #the values in the data in the order you want  
                  levels = c("NN", "JJ", "VB"),  
                  #give them better labels if you want  
                  labels = c("Noun", "Adjective", "Verb"))  
table(ELP$POS)
```

```
##
```

```
##      Noun Adjective      Verb
```

```
##      532       159       189
```


Dealing with Non-Normal Data

- ▶ One issue in language research is that often we have non-normal data
- ▶ Especially when working with frequency (as it is distributed by Zipf's law)

Dealing with Non-Normal Data

```
hist(ELP$SUBTLWF, breaks = 100)
```

Histogram of ELP\$SUBTLWF




Dealing with Non-Normal Data

- ▶ The simplest solution is to take the \log of the variable.
- ▶ Does make interpretation a bit more difficult, but helps with the distribution of the data.

Dealing with Non-Normal Data

```
ELP$Log_SUB = log(ELP$SUBTLWF)  
hist(ELP$Log_SUB)
```

Histogram of ELP\$Log_SUB



Build the Linear Model

- ▶ To be able to use our output for several purposes, we want to save it
 - ▶ You can call it whatever you want, I like `model`.
- ▶ Format for `lm` function is:
 - ▶ $Y \sim X + X + \dots$
 - ▶ `data =` name of data frame

```
model = lm(Mean_RT ~ Length + Log_SUB + POS,  
           data = ELP)
```

Summarize the Linear Model

```
summary(model)
```

```
##
```

```
## Call:
```

```
## lm(formula = Mean_RT ~ Length + Log_SUB + POS, data = EI)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

##	-213.70	-62.55	-9.71	53.87	389.00
----	---------	--------	-------	-------	--------

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

## (Intercept)	616.351	12.233	50.385	< 2e-16 ***
## Length	19.555	1.433	13.645	< 2e-16 ***
## Log_SUB	-29.288	1.784	-16.420	< 2e-16 ***
## POSAdjective	6.115	8.506	0.719	0.47238
## POSVerb	-23.069	7.918	-2.913	0.00367 **
## ---				

```
## Simf... 0.001... 0.001... 0.001... 0.001... 0.001...
```


Residuals

- ▶ A summary of the residuals - remember that residuals are the error terms or how far off we were at predicting the Mean_RT
- ▶ We will use this information as part of our assumptions diagnostics for data screening

```
summary(model$residuals)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	-213.699	-62.551	-9.714	0.000	53.874	388.998

Coefficients

- ▶ The coefficients table shows you the individual predictor significance levels
- ▶ If you use $p < .05$ as a criterion, we see that:
 - ▶ Intercept is the average RL
 - ▶ Length is a positive predictor: long words take longer to react to
 - ▶ Frequency is a negative predictor: more frequent words are faster (i.e. low freq = high RL)
 - ▶ Adjectives and Nouns have the same RL
 - ▶ Verbs and Nouns have different RL

Coefficients

```
options(scipen = 999)
round(summary(model)$coefficients, 3)
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	616.351	12.233	50.385	0.000
## Length	19.555	1.433	13.645	0.000
## Log_SUB	-29.288	1.784	-16.420	0.000
## POSAdjective	6.115	8.506	0.719	0.472
## POSVerb	-23.069	7.918	-2.913	0.004

Coefficients

- ▶ To interpret categorical predictors, it can help to make a means table
- ▶ Now I can see that verbs are responded to faster than nouns, interpreting the categorical predictor

```
tapply(ELP$Mean_RT, #dv  
       ELP$POS, #iv group variable  
       mean) #function
```

##	Noun	Adjective	Verb
##	787.5959	822.9145	754.3316

Coefficient Confidence Intervals

- We can calculate the confidence intervals for the coefficients, to help understand their precision

```
confint(model)
```

##		2.5 %	97.5 %
## (Intercept)	592.34193	640.36007	
## Length	16.74194	22.36757	
## Log_SUB	-32.78872	-25.78704	
## POSAdjective	-10.57915	22.80935	
## POSVerb	-38.61021	-7.52737	

Coefficient Practical Importance

- Calculate pr^2 : variance accounted for in the DV by that IV after removing all variance due to other IVs

$$\frac{t_{x_i}}{\sqrt{t_{x_i}^2 + df_{res}}}$$

```
t = summary(model)$coefficients[-1 , 3]
pr = t / sqrt(t^2 + model$df.residual)
pr^2
```

```
##           Length      Log_SUB POSAdjective      POSVerb
## 0.1754422571 0.2355448602 0.0005903474 0.0096065265
```

Overall Model

- ▶ So how much better than a random guess are we at predicting?
 - ▶ A good random guess is always the y-intercept or the mean of y.
- ▶ The F -statistic represents the difference of the model from zero

```
summary(model)$fstatistic
```

```
##      value      numdf      dendif  
## 183.7024    4.0000  875.0000
```

Overall Model

- ▶ R^2 represents the overlap in all IV variance with the DV variance

```
summary(model)$r.squared
```

```
## [1] 0.4564574
```


Diagnostic Tests


- ▶ Outliers and influential observations: data points that have large residuals or are otherwise odd in relation to the rest of the data
- ▶ Assumptions of parametric regression:
 - ▶ Independence
 - ▶ DV is response scale
 - ▶ Additivity (no multicollinearity)
 - ▶ Linearity
 - ▶ Normality
 - ▶ Homoscedasticity/Homogeneity

Outliers

- ▶ Hat values (or leverage): indicates how much influence on the slope a data point has
- ▶ Studentized residuals: the normalized (z-scored) difference between a participant's predicted and actual score
- ▶ Cook's values: a measure of influence (both leverage and discrepancy)

Outliers

```
library(car)  
influencePlot(model)
```



##

StudRes

Hat

CookD

Outliers

- What do we do with them?

```
ELP[c(331,660,498,411), ]
```

##		Word	Length	SUBTLWF	POS	Mean_RT	
## 331	interdepartmental		17	0.04	Adjective	1324.57	-
## 660	sacrilegious		12	0.39	Adjective	1228.06	-
## 498	whippet		7	0.10	Noun	1209.67	-
## 411	archenemy		9	0.25	Noun	1188.91	-

Assumptions

- ▶ Independence - the data is independent from each other (i.e. each data point is from a different “person”)
- ▶ Interval scale dependent variable: check!

Additivity

- ▶ No correlation between predictors above .9 (but .7 is actually not good either)

```
summary(model, correlation = T)$correlation[ , -1]
```

##	Length	Log_SUB	POSAdjective	
## (Intercept)	-0.94238493	-0.272940500	-0.09904041	-0.23
## Length	1.00000000	0.340416664	-0.05527619	0.06
## Log_SUB	0.34041666	1.000000000	0.09363100	0.00
## POSAdjective	-0.05527619	0.093631001	1.00000000	0.23
## POSVerb	0.06668845	0.006852748	0.23717736	1.00

Additivity


- ▶ Small Variance Inflation Scores (VIF) values (less than 5 to 10)

```
vif(model)
```

##		GVIF	Df	$GVIF^{(1/(2*Df))}$
##	Length	1.151054	1	1.072872
##	Log_SUB	1.150140	1	1.072446
##	POS	1.026925	2	1.006664

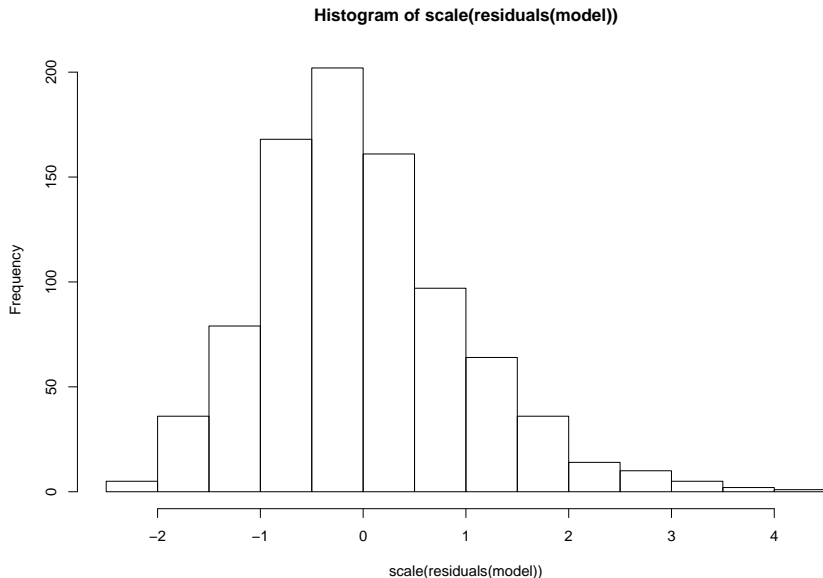
Linearity

```
plot(model, which = 2)
```




Normality

```
hist(scale(residuals(model)))
```




Homoscedasticity/Homogeneity

```
plot(model, which = 1)
```



Homoscedasticity/Homogeneity

```
{plot(scale(residuals(model)), scale(model$fitted.values))  
  abline(v = 0, h = 0)}
```



One Solution to Bad Assumptions

- First, build a function that saves the numbers you want:

```
bootcoef = function(formula, data, indices){  
  d = data[indices, ] #randomize the data by row  
  model = lm(formula, data = d) #run our model  
  return(coef(model)) #give back coefficients  
}
```

Bootstrapping

- ▶ Next, use the boot library to run the bootstraps (lots of runs on randomly sampled data)

```
library(boot)
model.boot = boot(formula = Mean_RT ~ Length + Log_SUB + PC,
                  data = ELP,
                  statistic = bootcoef,
                  R = 1000)
```

Bootstrapping

```
model.boot
```

```
##
```

```
## ORDINARY NONPARAMETRIC BOOTSTRAP
```

```
##
```

```
##
```

```
## Call:
```

```
## boot(data = ELP, statistic = bootcoef, R = 1000, formula
```

```
##       Length + Log_SUB + POS)
```

```
##
```

```
##
```

```
## Bootstrap Statistics :
```

```
##          original          bias      std. error
```

```
## t1* 616.350998  0.048541832    12.029762
```

```
## t2*  19.554757 -0.002752349     1.455837
```

```
## t3* -29.287879 -0.024054074     1.706049
```

```
## t4*   6.115101 -0.110164071     9.098897
```

```
## t5* -23.068792 -0.285394053     7.237738
```

CI for Bootstrapped Estimates

```
boot.ci(model.boot, index = 2)
```

```
## Warning in boot.ci(model.boot, index = 2): bootstrap var
```

```
## studentized intervals
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
```

```
## Based on 1000 bootstrap replicates
```

```
##
```

```
## CALL :
```

```
## boot.ci(boot.out = model.boot, index = 2)
```

```
##
```

```
## Intervals :
```

```
## Level      Normal                      Basic
```

```
## 95%    (16.70, 22.41 )    (16.63, 22.35 )
```

```
##
```

```
## Level      Percentile                  BCa
```

```
## 95%    (16.76, 22.48 )    (16.80, 22.60 )
```

```
## Calculations and Intervals on Original Scale
```