

## Lesson 9 - Topics Models

Erin M. Buchanan

02/28/2019

# Language Topics Discussed

- ▶ Expansion of semantic vector models into Topics Models
- ▶ Types of relations
- ▶ How to differentiate topics and other models

# Topics Background

- ▶ What does it take to understand a sentence?
  - ▶ Retrieving concepts from memory
  - ▶ Dynamic process based on incoming information
  - ▶ Use the semantic context to create a “gist” representation

# Topics Background


- ▶ Pulling the right information from memory can be improved by predicting what concepts are going to be relevant (expectancy generation)
  - ▶ For example, *bank* might prime *federal* and *reserve*
  - ▶ However, multiple senses can sometimes make this difficult
  - ▶ Gist representation allows us to create an overarching topic to disambiguate sense

# Topics Background

- ▶ Four types of ways to think about relation:
  - ▶ Word-concept: knowledge that a word refers to some concept (physical letters dog refer to dog)
  - ▶ Concept-concept: knowledge that a concept is related to some other concept (dog is a type of animal)
  - ▶ Concept-precept/action: knowledge about what a concept looks like or does (dogs are furry and bark)
  - ▶ Word-word: knowledge that the word co-occurs with another word (dog-cat)

# Topics Background

(a)



(b)

F  
FEL  
BA  
LOAN  
COMM  
STREAM  
RIVE  
DEEP  
MEADOW  
WOODS G

Figure 1 Approaches to semantic representation

# Topics Background

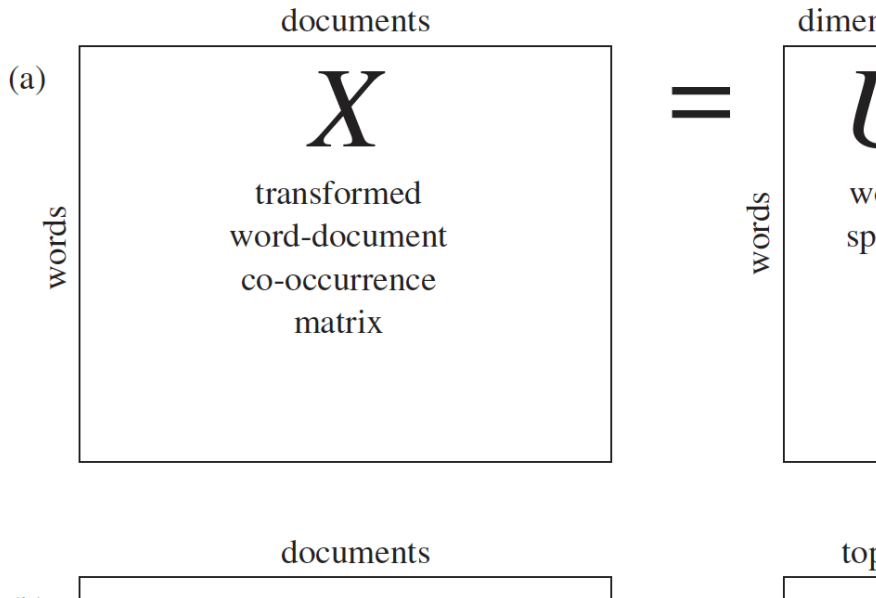
- ▶ These are useful to understand, because they predict different ways to think about semantic memory.
- ▶ What are people doing when they read a sentence and how can we represent that?
  - ▶ Translating words to concepts and using background knowledge to pull in other related concepts
  - ▶ Using word co-occurrence to predict the next words

# Topics Background

- ▶ What do topics models propose people do?
  - ▶ Predict: people predict the next word or concept because it helps with retrieval
  - ▶ Disambiguation: of senses or meanings of words
  - ▶ Gist: creating a coherent representation of the text (mental model not individual words)



## Topics Background



## Other thoughts

- ▶ Topics modeling could allow us to reveal topics present in text
- ▶ OR find ways to sort various texts into different groups
- ▶ Similar to clustering, classification, finds the natural groups in the corpus

## A little bit of math

- ▶ Latent Dirichlet allocation (LDA) is the most popular math
- ▶ Estimates topics based on the idea that every document includes a mix of topics, and every topic includes a mix of words
- ▶ That specification allows topics to overlap, such that they might have some of the same words/content
- ▶ LDA is the middle group that finds both the words for each topic and the topics for each document

# Getting started with raw data

- First, you would need to load the libraries for the Topic Modeling packages:

```
library(tm)
library(topicmodels)
library(tidyverse)
library(tidytext)
library(slam)
```

## Load a dataset or corpus

- ▶ Then, you could load a dataset you are interested in working with:

```
import df = read.csv('exam_answers.csv', header = F, stringsAsFactors = F)
```

## Convert to a Corpus

- ▶ From these documents, we will create a corpus (a set of text documents).
- ▶ Because our data is in one column in our dataset, we will use `VectorSource()` to create the corpus:

```
import_corpus = Corpus(VectorSource(importdf$V1))
```

## Clean up the text

- ▶ When you perform these analyses, you usually have to edit the text.
- ▶ Therefore, we are going to lower case the words, take out the punctuation, and remove English stop words (like *the*, *an*, *a*).
- ▶ This step will also transform the documents in a term (words) by document matrix.

## Clean up the text

```
import_mat =  
    DocumentTermMatrix(import_corpus,  
        control = list(stemming = TRUE, #create root words  
                        stopwords = TRUE, #remove stop words  
                        minWordLength = 3, #cut out small words  
                        removeNumbers = TRUE, #take out numbers  
                        removePunctuation = TRUE)) #take out punctuation
```



## Weight the matrix

- ▶ Then you would want to weight that matrix to help control for the sparsity of the matrix.
- ▶ That means you are controlling for the fact that not all words are in each document, as well as the fact that some words are very frequent.
- ▶ Then you usually ignore very frequent words and words with zero frequency.

## Weight the matrix

*#weight the space*

```
import_weight = tapply(import_mat$v/row_sums(import_mat)[import_mat$term,
                                             import_mat$j,
                                             mean) *
  log2(nDocs(import_mat)/col_sums(import_mat > 0))
```

*#ignore very frequent and 0 terms*

```
import_mat = import_mat[ , import_weight >= .1]
import_mat = import_mat[ row_sums(import_mat) > 0, ]
```

## Parameters and terms

- ▶ Alpha: a measure of the number of topics; low scores indicate a few dominant topics per document, high scores indicate more
- ▶ Beta: a measure of the number of words, low scores indicate each topic only composes of a few words
- ▶ Gamma: probability of that topic in that document
- ▶ Entropy: a measure of randomness

## A bit more math

- ▶ There are several model types:
  - ▶ The LDA Fit model is an analysis with VEM (variational expectation-maximization) algorithm and estimating an alpha.
    - ▶ The LDA Fixed model using the VEM algorithm with a fixed alpha value.
    - ▶ Last, the LDA Gibbs option uses a Gibbs (Bayesian) algorithm to fit the data.
    - ▶ CTM stands for correlated topics models, which allows the correlation between topics, and this method uses a VEM algorithm.

# Run the models

- ▶ First, you will pick a number of expected topics - which is the  $k$  option.
- ▶ The SEED should be a random number to start the analysis on.

```
k = 3 #set the number of topics
```

```
SEED = 2010 #set a random number
```

```
LDA_fit = LDA(import_mat, k = k,  
               control = list(seed = SEED))
```

```
LDA_fixed = LDA(import_mat, k = k,  
                 control = list(estimate.alpha = FALSE, seed = SEED))
```

## Run the models

```
LDA_gibbs = LDA(import_mat, k = k, method = "Gibbs",
                 control = list(seed = SEED, burnin = 1000,
                               thin = 100, iter = 1000))

CTM_fit = CTM(import_mat, k = k,
               control = list(seed = SEED,
                             var = list(tol = 10^-4),
                             em = list(tol = 10^-3)))
```

## Get the alpha values

- ▶ You can then get the alpha values, and smaller alpha values indicate higher percentages of documents that were classified to one single topic.

```
LDA_fit@alpha
```

```
## [1] 0.05846617
```

```
LDA_fixed@alpha
```

```
## [1] 16.66667
```

```
LDA_gibbs@alpha
```

```
## [1] 16.66667
```

## Get the entropy values

- ▶ You can also get entropy values where higher values indicate that topics are evenly spread.

```
sapply(list(LDA_fit, LDA_fixed, LDA_gibbs, CTM_fit),  
       function (x)  
         mean(apply(posterior(x)$topics, 1, function(z) - s
```

```
## [1] 0.2402260 1.0920326 1.0946644 0.5894953
```



## The actual topics

- ▶ The topic matrix indicates the rank of the number of topics for each document.
- ▶ For instance, if you select to estimate 5 topics, you will see which topic is covered most in each document, with less covered topics ranked lower.
- ▶ Therefore, a score set of 5, 3, 1, 2, 4 indicates that the 5th topic was covered most in that document, and the 4th topic was covered least.

## The actual topics

```
topics(LDA_fit, k)
```

```
##          1 2 3 4 5 6 7 8 9 10 11 14 15 16 17 18 19 20 21 22
## [1,] 3 3 1 3 3 1 1 1 2  1  1  2  3  2  3  1  3  3  1  2
## [2,] 2 1 3 1 1 2 3 2 1  3  2  1  1  1  2  2  1  1  2  1
## [3,] 1 2 2 2 2 3 2 3 3  2  3  3  2  3  1  3  2  2  3  3
##          30 31 32 33 34 35 36 37 38 39 40 41 42
## [1,]  1  1  1  3  2  2  3  2  2  1  2  2  2
## [2,]  2  3  3  1  3  1  1  1  1  2  1  1  3
## [3,]  3  2  2  2  1  3  2  3  3  3  3  3  1
```

## The actual topics

```
topics(LDA_gibbs, k)
```

```
##           1 2 3 4 5 6 7 8 9 10 11 14 15 16 17 18 19 20 21 22
## [1,] 2 3 1 2 3 1 1 2 2 2 2 2 2 1 1 1 3 1 2 2
## [2,] 1 1 2 1 1 2 2 1 1 1 3 1 1 2 2 2 1 2 3 1
## [3,] 3 2 3 3 2 3 3 3 3 3 1 3 3 3 3 3 2 3 1 3
##           30 31 32 33 34 35 36 37 38 39 40 41 42
## [1,] 1 3 1 1 2 3 2 2 2 3 3 1 2
## [2,] 2 1 2 3 1 1 1 1 3 1 1 2 1
## [3,] 3 2 3 2 3 2 3 3 1 2 2 3 3
```

*##you can do all of them saving space*

*#topics(LDA\_fixed, k)*

*#topics(CTM\_fit, k)*

## The terms for topics

- You can get the most frequent terms for each of the topics that were estimated.

```
terms(LDA_fit,10)
```

##		Topic 1	Topic 2	Topic 3
##	[1,]	"top"	"element"	"bike"
##	[2,]	"actual"	"observ"	"mean"
##	[3,]	"previous"	"gorilla"	"success"
##	[4,]	"surround"	"black"	"ride"
##	[5,]	"environ"	"line"	"gorilla"
##	[6,]	"dont"	"screen"	"red"
##	[7,]	"discuss"	"demonstr"	"dont"
##	[8,]	"therefor"	"flash"	"your"
##	[9,]	"hand"	"lead"	"environ"
##	[10,]	"rememb"	"monitor"	"effici"

# The terms for topics

```
terms(LDA_gibbs,10)
```

##		Topic 1	Topic 2	Topic 3
##	[1,]	"mean"	"gorilla"	"bike"
##	[2,]	"dont"	"observ"	"success"
##	[3,]	"environ"	"surround"	"ride"
##	[4,]	"element"	"hand"	"actual"
##	[5,]	"top"	"therefor"	"previous"
##	[6,]	"rememb"	"your"	"red"
##	[7,]	"screen"	"attentioncontrol"	"black"
##	[8,]	"act"	"demonstr"	"discuss"
##	[9,]	"although"	"lead"	"line"
##	[10,]	"close"	"neutral"	"monitor"

```
##again you can do all of them  
#terms(LDA_fixed, 10)  
#terms(CTM_fit,10)
```

## Make some pretty plots

```
#use tidyverse to clean up the the fit  
LDA_fit_topics = tidy(LDA_fit, matrix = "beta")  
  
#create a top terms  
top_terms = LDA_fit_topics %>%  
  group_by(topic) %>%  
  top_n(10, beta) %>%  
  ungroup() %>%  
  arrange(topic, -beta)
```

# Make some pretty plots

- Some code to clean up the ggplot2 defaults

```
cleanup = theme(panel.grid.major = element_blank(),  
                panel.grid.minor = element_blank(),  
                panel.background = element_blank(),  
                axis.line.x = element_line(color = "black"),  
                axis.line.y = element_line(color = "black"),  
                legend.key = element_rect(fill = "white"),  
                text = element_text(size = 10))
```

# Make some pretty plots

```
#make the plot
```

```
top_terms %>%
```

```
  mutate(term = reorder(term, beta)) %>%
```


```
  ggplot(aes(term, beta, fill = factor(topic))) +
```

```
  geom_bar(stat = "identity", show.legend = FALSE) +
```

```
  facet_wrap(~ topic, scales = "free") +
```


```
  cleanup +
```

```
  coord_flip()
```





# The plot




# Document classification

- ▶ We saw earlier with the `topics()` function, we could figure out the most to least likely topics.
- ▶ This matrix is organized by gamma, which is the probability of that topic in for each document.
- ▶ Let's visualize those numbers.


# Document classification

```
LDA_gamma = tidy(LDA_fit, matrix = "gamma")
```

```
LDA_gamma %>%  
  ggplot(aes(factor(topic), gamma)) +  
  geom_point() +  
  cleanup
```



# Document classification



# Summary

- ▶ We explored how the theoretical background for topics models is different than other semantic vector space models.
- ▶ We talked about how to build topics models with various settings.
- ▶ We talked about the output you can pull from a topic model.
- ▶ Extensions can be made to unsupervised classification and clustering.